

# Announcements

- Assignment 1 has been posted on **Crowdmark**.  
You can get access to **Crowdmark** through **myCourses**.
- TAs' Zoom links have been posted on myCourses.

TAs will start office hours next week.

Two TAs are responsible for each week's office hours and each assignment.

# Computer Numbers and Arithmetic


- 3 lectures
- References:
  - Overton, Numerical Computing with IEEE Floating Point Arithmetic, SIAM, 2004.
  - IEEE Computer Society, IEEE Standard for Floating-Point Arithmetic, 2019.
  - Cheney & Kincaid, Numerical Mathematics and Computing, Sections 1.1 and 1.3.

# Today's topics

- Binary and decimal representations of real numbers (Overton Chap 2).
- Computer representation of integers and non-integral real numbers (Overton Chap 3).
- IEEE floating point representation (Overton Chap 4)

# Classes of Real Numbers

All real numbers can be represented by a line:



The Real Line

$$\text{real numbers} \left\{ \begin{array}{l} \text{rational numbers} \left\{ \begin{array}{l} \text{integers} \\ \text{nonintegral fractions} \end{array} \right. \\ \text{irrational numbers} \end{array} \right.$$

# Classes of Real Numbers

- **Rational numbers:**

all the real numbers which consist of a ratio of two integers.  
e.g.,  $2/1, 1/3, \dots$

- **Irrational numbers:**

Most real numbers are **not** rational, i.e. there is no way of writing them as the ratio of two integers.

Familiar examples of irrational numbers are:

$$\sqrt{2}, \quad \pi, \quad e \equiv \lim_{n \rightarrow \infty} \left(1 + \frac{1}{n}\right)^n.$$

## How to represent numbers?

Two basic systems:

- The **decimal**, or **base 10**, system requires 10 symbols, 0, 1, 2, 3, 4, 5, 6, 7, 8, 9.
- The **binary**, or **base 2**, system is convenient for electronic computers.  
Every number is represented as a string of **0**'s and **1**'s.

## Decimal & binary representations (expansions)

- **Integers:**

The decimal and binary representation of **integers** requires an expansion in nonnegative powers of the base; e.g.

$$(71)_{10} = 7 \times 10 + 1$$

its **binary equivalent:**  $(1000111)_2 =$

$$1 \times 2^6 + 0 \times 2^5 + 0 \times 2^4 + 0 \times 2^3 + 1 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$$

- **Non-integral rational numbers:**

They have entries to the **right** of the decimal or binary point

## Representations of nonintegral rational numbers

- Both representations are **finite**:

$$\frac{11}{2} = (5.5)_{10} = 5 \times 1 + 5 \times \frac{1}{10},$$

$$\frac{11}{2} = (101.1)_2 = 1 \times 4 + 0 \times 2 + 1 \times 1 + 1 \times \frac{1}{2}$$

- The decimal is finite, but the **binary is infinitely long**

$$1/10 = (0.1)_{10}$$

$$\begin{aligned} \frac{1}{10} &= (0.000\textcolor{red}{1}\textcolor{red}{0}\textcolor{red}{0}\textcolor{blue}{1}\textcolor{blue}{0}\textcolor{blue}{0}\textcolor{blue}{1}\textcolor{blue}{0}\textcolor{blue}{0} \dots)_2 \quad (\text{it is repeating}). \\ &= \frac{1}{16} + \frac{1}{32} + \frac{0}{64} + \frac{0}{128} + \frac{1}{256} + \frac{1}{512} + \dots \end{aligned}$$



## Representations of nonintegral rational numbers, ctd

- Both representations are **infinite** and **repeating**:

$$1/3 = (0.333\dots)_{10} = (0.010101\dots)_2.$$

If the representation of a **rational** number is *infinite*, it **must** be *repeating*.

- Is it possible that the decimal representation is infinite, but the binary representation is finite?

NO

## Representations of irrational numbers

- **Irrational** numbers:

**Irrational** numbers always have **infinite, non-repeating** expansions. e.g.

$$\sqrt{2} = (1.414213\dots)_{10},$$

$$\pi = (3.141592\dots)_{10},$$

$$e = (2.71828182845\dots)_{10}.$$

## Conversion between binary & decimal

- **Binary**  $\rightarrow$  **decimal**:

Easy. e.g.  $(1001.11)_2$  is the decimal number

$$1 \times 2^3 + 0 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 + 1 \times 2^{-1} + 1 \times 2^{-2} = 9.75$$

- **Decimal**  $\rightarrow$  **binary**:

Convert the integer and fractional parts separately.

e.g. if  $x$  is a **decimal integer**, we want to find  $a_0, a_1, \dots, a_n$ , all 0 or 1 such that  $(x)_{10} = (a_n a_{n-1} \dots a_0)_2$ , i.e.,

$$a_n \times 2^n + a_{n-1} \times 2^{n-1} + \dots + a_1 \times 2 + a_0 \times 2^0 = x,$$

Clearly dividing  $x$  by 2 gives **remainder**  $a_0$ , leaving as **quotient**

$$a_n \times 2^{n-1} + a_{n-1} \times 2^{n-2} + \dots + a_1 \times 2^0,$$

and so we can continue to find  $a_1$  then  $a_2$  etc.

## Converting between binary & decimal

**Q:** What is a similar approach for decimal fractions?

# Computer Representation of Numbers

- Integers

- ① sign-and-modulus
- ② 2's complement representation

- Non-integral real numbers

- ① Fixed point
- ② Floating point

## Computer representation of integers

Typically, integers are stored using a 32-bit word.

1. **sign-and-modulus** — a simple approach.

Use 1 bit to represent the **sign**, and store the **binary** representation of the magnitude of the integer. e.g. decimal 71 is stored as the bitstring

0	00...01000111
---	---------------

**Q.** What is the **largest** magnitude which fits a 32-bit word?

0	111111...1111
---	---------------

The largest magnitude is  $2^{31} - 1$

## Explanation of Ariane 5 Explosion (1996)

The rocket required the conversion of a 64-bit floating point number to a 16 bit signed integer.

The largest integer number which can fit a 16-bit word is

$$2^{15} - 1 = 32,767.$$

But the number was larger than 32,767.

The conversion failed.

## Computer representation of integers

### 2. 2's complement representation (CR)

– more convenient & used by most machines.

- (i) The **nonnegative** integers 0 to  $2^{31} - 1$  are stored as before, e.g., 71 is stored as the bitstring 000...01000111
- (ii) A **negative integer**  $-x$ , where  $1 \leq x \leq 2^{31}$ , is stored as the **positive integer**  $2^{32} - x$ .

e.g.  $-71$  is stored as the bitstring 111...10111001.

**Converting  $x$  to the 2's CR  $2^{32} - x$  of  $-x$ :**

$$\begin{aligned} 2^{32} - x &= (2^{32} - 1 - x) + 1, \\ 2^{32} - 1 &= (111 \dots 111)_2. \end{aligned}$$

**Chang all 0 bits of  $x$  to 1s, 1 bits to 0 and add 1.**



## 2's complement representation (CR)

**Q:** *What is the quickest way of deciding if a number is negative or nonnegative using 2's CR ?*

**A:** Check the leading bit.

If it's 1, the number is negative, else the number is nonnegative.

## 2's complement representation — an advantage

Consider  $y + (-x)$ , where  $1 \leq x \leq 2^{31}$ ,  $0 \leq y \leq 2^{31} - 1$ .

2's CR of  $y$  is  $y$ ;    2's CR of  $-x$  is  $2^{32} - x$

**Adding** these two **representations** gives

$$y + (2^{32} - x) = 2^{32} + y - x = 2^{32} - (x - y).$$

- If  $y \geq x$ , the LHS will not fit in a 32-bit word, and the **leading bit** can be dropped, giving the **correct result**,  $y - x$ .
- If  $y < x$ , the RHS is **already correct**, since it *represents*  $-(x - y)$ .

Thus, **no special hardware is needed for integer subtraction**.

The addition hardware can be used, once  $-x$  has been represented using 2's complement.

## Computer representation of nonintegral real numbers

**Real** numbers are approximately stored using the **binary** representation of the number.

Two possible methods: **fixed point** and **floating point**.

### Fixed point:

the computer word is divided into **three fields**, one for each of:

- the **sign** of the number
- the number **before** the point
- the number **after** the point.

In a **32-bit word** with field widths of **1,15** and **16**, the number  $11/2$  would be stored as:

0	0000000000000101	1000000000000000
---	------------------	------------------

## Fixed point system

The fixed point system has a severe limitation on the **size** of the numbers to be stored.

Smallest magnitude:

$$\boxed{0 \mid 0000000000000000 \mid 0000000000000001} = 2^{-16}$$

Largest magnitude:

$$\boxed{0 \mid 1111111111111111 \mid 1111111111111111} = 2^{15} - 2^{-16}$$

A fix point system **is inadequate for most scientific computing**.

But it is fast and is used in some real-time applications.

# Explanation of the Patriot Missile Failure (1991)

- The missile system's internal clock stored the time as an integer value in **units of tenths of a second**.
- To produce the time in seconds, multiply the stored integer by  $1/10$ .
- The binary expansion of  $1/10$  is  
 $0.00011001100110011001100|11001100\dots$
- The system stored the 23 bits after the binary point  
 $0.00011001100110011001100$
- The chopping error is  $(0.11001100\dots)_2 \times 2^{-23} \approx 9.5 \times 10^{-8}$
- After 100 hours, the stored integer is  $100 \times 60 \times 60 \times 10$ .
- The error in time after 100 hours:

$$100 \times 60 \times 60 \times 10 \times 9.5 \times 10^{-8} \approx 0.34 \text{ seconds}$$

# IEEE Floating Point Standard

## Motivations:

- Floating point computation was in standard use by the mid 1950s.
- During the subsequent two decades, each computer manufacturer developed its own floating point system, leading to much inconsistency in how one program might behave on different machines.
- It was very difficult to write **portable software** that would work properly on all machines.

## IEEE Floating Point Standard

- **IEEE 754-1985**: a **binary** floating point standard. It was developed through the efforts of **William Kahan** & others.
- **IEEE 854-1987**: radix independent floating point arithmetic. It was motivated by decimal (radix-10) machines, and set the requirement for both binary and decimal floating point arithmetic in a common framework,
- **IEEE 754-2008**: a significant revision of IEEE 754-1985. It extended the previous standard where it was necessary, added decimal arithmetic and formats, and merged in IEEE 854-1987.
- **IEEE 754-2019**: a minor revision of IEEE 754-2008, incorporating mainly clarifications, defect fixes and new recommended operations.

In this course, the “IEEE standard” refers to the binary standard.

# IEEE Floating Point Standard

The standard defines:

- **arithmetic formats:** floating-point formats that can be used to represent floating-point operands or results for the operations of the standard.
- **interchange formats:** formats that have a specific fixed-width encodings (bit strings) that may be used to exchange floating-point data in an efficient and compact form
- **rounding rules:** properties to be satisfied when rounding numbers during arithmetic and conversions
- **operations:** arithmetic and other operations (such as trigonometric functions) on operands
- **exception handling:** indications of exceptional conditions (such as division by zero, overflow, etc.)



# IEEE Floating Point Formats

The IEEE standard has 3 binary floating point basic formats :

binary32, single format, or single precision;

binary64, double format, or double precision;

binary128, quadruple format, or quadruple precision

They have different numbers of bits to represent the significand and exponent.

## Base-2 Normalized Exponential Notation

In base-2 normalized exponential notation, any nonzero binary number  $x$  can be written as

$$x = (-1)^s \times m \times 2^E, \quad \text{where } 1 \leq m < 2,$$

where

- $s$  is 0 or 1.
- $m$  is called the **significand** or **mantissa**, and its binary expansion is


$$m = (b_0.b_1b_2b_3\dots)_2, \quad \text{with } b_0 = 1.$$

- $E$  is an **integer**, called the **exponent**.

## IEEE Single Format (binary32)

$$x = (-1)^s \times m \times 2^E, \quad m = (b_0.b_1b_2b_3\dots)_2, \quad b_0 = 1,$$

**Single format** numbers use **32-bit words**.



- **sign field:** 1 bit for  $s$ .
- **exponent field:** 8 bits  $a_1a_2a_3\dots a_8$  for  $E$ ,  $E_{\min} \leq E \leq E_{\max}$ .
- **significand field:** 23 bits for  $m$ .

$b_0$  is not stored, since it is known that  $b_0 = 1$ .

This idea is called **hidden bit normalization**.

The significand field is also referred to as the **fraction field**.

## IEEE Single Format, ctd

It may not be possible to store  $x = \pm(b_0.b_1b_2b_3\dots)_2 \times 2^E$  exactly with such a scheme, because

- either  $E$  is outside the permissible range.
- or  $b_{24}, b_{25}, \dots$  are **not all zero**.

**Def.** A number is called a **(computer) floating point number** if it can be stored **exactly** this way, e.g.,

$$71 = (1.000111)_2 \times 2^6$$

can be represented by

0	$E = 6$	0001110000000000000000
---	---------	------------------------

If  $x$  is not a floating point number, it must be **rounded** before it can be stored on the computer.

## Special Numbers

- 0. Zero cannot be **normalized**.
- $-0$ .  $-0$  and  $0$  are **two different representations for the same number**
- $\infty$ . This allows e.g.  $1.0/0.0 \rightarrow \infty$ , instead of terminating with an **overflow** message.
- $-\infty$ .  $-\infty$  and  $\infty$  represent **two very different numbers**.
- NaN, or “**Not a Number**”, and is an **error pattern**.
- Subnormal numbers (see later)

All special numbers are represented by a **special bit pattern** in the **exponent field**.

## IEEE Single Format (binary32)

$\pm$	$a_1 a_2 a_3 \dots a_8$	$b_1 b_2 b_3 \dots b_{23}$
-------	-------------------------	----------------------------

If exponent $a_1 \dots a_8$ is	Then value is
$(00000000)_2 = (0)_{10}$	$\pm(0.b_1..b_{23})_2 \times 2^{-126}$
$(00000001)_2 = (1)_{10}$	$\pm(1.b_1..b_{23})_2 \times 2^{-126}$
$(00000010)_2 = (2)_{10}$	$\pm(1.b_1..b_{23})_2 \times 2^{-125}$
$\downarrow$	$\downarrow$
$(01111111)_2 = (127)_{10}$	$\pm(1.b_1..b_{23})_2 \times 2^0$
$(10000000)_2 = (128)_{10}$	$\pm(1.b_1..b_{23})_2 \times 2^1$
$\downarrow$	$\downarrow$
$(11111101)_2 = (253)_{10}$	$\pm(1.b_1..b_{23})_2 \times 2^{126}$
$(11111110)_2 = (254)_{10}$	$\pm(1.b_1..b_{23})_2 \times 2^{127}$
$(11111111)_2 = (255)_{10}$	$\pm\infty$ if $b_1, \dots, b_{23} = 0$ ; NaN otherwise.

## IEEE Single Format, ctd.

$\pm$	$a_1 a_2 a_3 \dots a_8$	$b_1 b_2 b_3 \dots b_{23}$
-------	-------------------------	----------------------------

All lines except the first and the last refer to the **normal** numbers, i.e. **not special**.

The exponent representation uses **biased representation**:  
this bitstring is the binary representation of  $E + 127$ .

127 is the **exponent bias**. e.g.  $1 = (1.000\dots 0)_2 \times 2^0$  is stored as

0	01111111	000000000000000000000000
---	----------	--------------------------

**Exponent range** for normal numbers is  
00000001 to 11111110 (1 to 254), representing **actual exponents**

$E_{\min} = -126 \text{ to } E_{\max} = 127$
--

## IEEE Single Format, ctd.

$\pm$	$a_1 a_2 a_3 \dots a_8$	$b_1 b_2 b_3 \dots b_{23}$
-------	-------------------------	----------------------------

The **smallest normal positive number** is

$$(1.000 \dots 0)_2 \times 2^{-126} :$$

0	00000001	000000000000000000000000
---	----------	--------------------------

approximately  $1.2 \times 10^{-38}$ .

The **largest normal positive number** is

$$(1.111 \dots 1)_2 \times 2^{127} :$$

0	11111110	111111111111111111111111
---	----------	--------------------------

approximately  $3.4 \times 10^{38}$ .



## IEEE Single Format, ctd.

$\pm$	$a_1 a_2 a_3 \dots a_8$	$b_1 b_2 b_3 \dots b_{23}$
-------	-------------------------	----------------------------

Last line:

If exponent $a_1 \dots a_8$ is	Then value is
$(11111111)_2 = (255)_{10}$	$\pm\infty$ if $b_1, \dots, b_{23} = 0$ ; NaN otherwise

This shows an **exponent bitstring of all ones** is a special pattern for  $\pm\infty$  or NaN, depending on the value of the fraction.

NaN has two types:

quiet NaN (qNaN) if  $b_1 = 1$ , signaling NaN (sNaN) if  $b_1 = 0$ .

## IEEE Single Format, ctd.

$\pm$	$a_1 a_2 a_3 \dots a_8$	$b_1 b_2 b_3 \dots b_{23}$
-------	-------------------------	----------------------------

First line

$(00..00)_2 = (0)_{10}$	$\pm(0.b_1..b_{23})_2 \times 2^{-126}$
-------------------------	--

**zero** requires a zero bitstring for the *exponent* field **as well as** for the *fraction*:

0	00000000	000000000000000000000000
---	----------	--------------------------

**Initial unstored bit is 0, not 1, in line 1.**

## IEEE Single Format, ctd.

*First line*

$$\boxed{(00..00)_2 = (0)_{10} \mid \pm(0.b_1..b_{23})_2 \times 2^{-126}}$$

If **exponent** is **zero**, but **fraction** is **nonzero**, the number represented is **subnormal**.

Although  $2^{-126}$  is the smallest normal positive number, we can represent smaller **subnormal** numbers.

e.g.  $2^{-127} = (0.1)_2 \times 2^{-126}$ .

0	00000000	10000000000000000000000000
---	----------	----------------------------

and  $2^{-149} = (0.0000 \dots 01)_2 \times 2^{-126}$ .

[illegible]

This is the smallest positive number we can store.

## IEEE Single Format, ctd.

Subnormal numbers cannot be normalized, as this would give exponents which do not fit.

Subnormal numbers are **less accurate** (less room for nonzero bits in the fraction). e.g.,

$$(1/10) \times 2^{-133} = (0.11001100\dots)_2 \times 2^{-136}$$

is

0	00000000	000000000001100110011001
---	----------	--------------------------

## IEEE Single Format, ctd.

**Q:** (i) How is 2 represented ??

0	10000000	000000000000000000000000
---	----------	--------------------------

(ii) What is the next smallest IEEE single precision number larger than 2 ??

0	10000000	000000000000000000000001
---	----------	--------------------------

(iii) What is the gap between 2 and the first IEEE single precision number larger than 2?

$$2^{-23} \times 2 = 2^{-22}.$$

### General Result:

Let  $x = m \times 2^E$  be a single format number. The **gap** between  $x$  and the next smallest single format number larger than  $x$  is

$$2^{-23} \times 2^E.$$

## Precision, Machine Epsilon

**Def. Precision** : The number of bits in the significand (including the hidden bit) is called the **precision** of the floating point system, denoted by  $p$ .

In the **single format** system,  $p = 24$ , the “single precision” is 24. Recall “single precision” also refers to the single format.

**Def. Machine Epsilon** : The gap between the number **1** and the **next larger** floating point number is called the **machine epsilon** of the floating point system, denoted by  $\epsilon$ .

In the **single format** system, the number after 1 is

$$b_0.b_1 \dots b_{23} = 1.000000000000000000000001,$$

so  $\epsilon = 2^{-23}$ .


## IEEE Double Format (binary64)

Each **double format** floating point number is stored in a **64-bit double word**.

Ideas are the same:

Field widths (1, 11 & 52) and exponent bias (1023) different.

$b_1, \dots, b_{52}$  can be stored instead of  $b_1, \dots, b_{23}$ .






# IEEE Double Format, ctd

If exponent is $a_1..a_{11}$	Then value is
$(000..0000)_2 = (0)_{10}$	$\pm(0.b_1..b_{52})_2 \times 2^{-1022}$
$(000..0001)_2 = (1)_{10}$	$\pm(1.b_1..b_{52})_2 \times 2^{-1022}$
$(000..0010)_2 = (2)_{10}$	$\pm(1.b_1..b_{52})_2 \times 2^{-1021}$
↓	↓
$(01..111)_2 = (1023)_{10}$	$\pm(1.b_1..b_{52})_2 \times 2^0$
$(10..000)_2 = (1024)_{10}$	$\pm(1.b_1..b_{52})_2 \times 2^1$
↓	↓
$(11..101)_2 = (2045)_{10}$	$\pm(1.b_1..b_{52})_2 \times 2^{1022}$
$(11..110)_2 = (2046)_{10}$	$\pm(1.b_1..b_{52})_2 \times 2^{1023}$
$(11..111)_2 = (2047)_{10}$	$\pm\infty$ if $b_1, , b_{52} = 0$ ; NaN otherwise

## IEEE Quadruple Format (binary128)

The **quadruple format (binary128)** uses 128 bits, the field widths are (1, 15 & 112) and exponent bias is  $2^{14} - 1 = 16383$ .



# IEEE Quadruple Format (binary128), cod

$\pm$	$a_1 a_2 a_3 \dots a_{15}$	$b_1 b_2 b_3 \dots b_{112}$
If exponent is $a_1..a_{15}$		Then value is
$(000..0000)_2 = (0)_{10}$		$\pm(0.b_1..b_{112})_2 \times 2^{-16382}$
$(000..0001)_2 = (1)_{10}$		$\pm(1.b_1..b_{112})_2 \times 2^{-16382}$
$(000..0010)_2 = (2)_{10}$		$\pm(1.b_1..b_{112})_2 \times 2^{-16381}$
$\downarrow$		$\downarrow$
$(01..111)_2 = (16383)_{10}$		$\pm(1.b_1..b_{112})_2 \times 2^0$
$(10..000)_2 = (16384)_{10}$		$\pm(1.b_1..b_{112})_2 \times 2^1$
$\downarrow$		$\downarrow$
$(11..101)_2 = (32765)_{10}$		$\pm(1.b_1..b_{112})_2 \times 2^{16382}$
$(11..110)_2 = (32766)_{10}$		$\pm(1.b_1..b_{112})_2 \times 2^{16383}$
$(11..111)_2 = (32767)_{10}$		$\pm\infty$ if $b_1, \dots, b_{112} = 0$ ; NaN otherwise

## Machine Precision, Epsilon of the 3 Formats

Single	$p = 24$	$\epsilon = 2^{-23} \approx 1.2 \times 10^{-7}$
Double	$p = 53$	$\epsilon = 2^{-52} \approx 2.2 \times 10^{-16}$
Quadruple	$p = 113$	$\epsilon = 2^{-112} \approx 1.9 \times 10^{-34}$

# Rounding

We use **Floating Point Data** to include

$\pm 0$ , subnormal & normal FPNs, &  $\pm\infty$  and NaNs

in a given format, e.g., single. These form a **finite** set.

$N_{\min}$ : the minimum positive **normal** FPN;

$N_{\max}$ : the maximum positive **normal** FPN;

A real number  $x$  is in the “**normal range**” if

$$N_{\min} \leq |x| \leq N_{\max}$$

## Rounding, ctd.

### Question:

Let  $x$  be a real number and  $|x| \leq N_{\max}$ .

If  $x$  is not a floating point number,

what are two obvious choices for the floating point **approximation** to  $x$  ??

$x_-$  : the closest FPN **less** than  $x$ ,

$x_+$  : the closest FPN **greater** than  $x$ .

## Rounding, ctd.

Consider the single format.

Let  $x$  be positive with

$$x = (b_0.b_1b_2 \dots b_{23}b_{24}b_{25} \dots)_2 \times 2^E$$

$b_0 = 1$  (normal), or  $b_0 = 0$ ,  $E = -126$  (subnormal).

Then **discard**  $b_{24}$ ,  $b_{25}$ , ... gives

$$x_- = (b_0.b_1b_2 \dots b_{23})_2 \times 2^E.$$

An algorithm for  $x_+$  is more complicated since it may involve some bit “carries”.

$$x_+ = [(b_0.b_1b_2 \dots b_{23})_2 + (0.00 \dots 1)_2] \times 2^E.$$

If  $x$  is **negative**, the situation is reversed:  $x_+$  is obtained by dropping bits  $b_{24}$ ,  $b_{25}$ , etc.

## Correctly Rounded Values

The IEEE standard defines the **correctly rounded value** of  $x$ ,  $\text{round}(x)$ .

If  $x$  is a floating point number,  $\text{round}(x) = x$ .

Otherwise  $\text{round}(x)$  depends on the **rounding mode** in effect:

- **Round down:**  $\text{round}(x) = x_-$ .
- **Round up:**  $\text{round}(x) = x_+$ .
- **Round towards zero:**  $\text{round}(x)$  is either  $x_-$  or  $x_+$ , whichever is between zero and  $x$ .
- **Round to nearest:**  $\text{round}(x)$  is either  $x_-$  or  $x_+$ , whichever is nearer to  $x$ . In the case of a tie, the one with its **least significant bit (the last bit) equal to zero** is chosen.

This rounding mode is almost always used.



## Correctly Rounded Values

If  $x$  is **positive**, then  $x_-$  is between zero and  $x$ ,  
so **round down** and **round towards zero** have the same effect;  
**round towards zero** simply requires **truncating** the binary  
expansion, i.e. discarding bits.

**“Round”** with no qualification usually means **“round to nearest”**.

## Absolute Rounding Error

**Def.** The **absolute rounding error** associated with  $x$ :

$$|\text{round}(x) - x|.$$

Its value depends on mode.

For all modes  $|\text{round}(x) - x| < |x_+ - x_-|$ .

Suppose  $N_{\min} \leq x \leq N_{\max}$ ,

$$x = (b_0.b_1b_2 \dots b_{23}b_{24}b_{25} \dots)_2 \times 2^E, \quad b_0 = 1$$

$$\text{IEEE single } x_- = (b_0.b_1b_2 \dots b_{23})_2 \times 2^E$$

$$\text{IEEE single } x_+ = x_- + (0.00 \dots 01)_2 \times 2^E$$

So for **any** mode

$$|\text{round}(x) - x| < 2^{-23} \times 2^E$$

## Absolute Rounding Error

In general for **any rounding mode**:

$$|\text{round}(x) - x| < \epsilon \times 2^E \quad (*)$$

**Q:** Does  $(*)$  hold if  $0 < x < N_{\min}$ , i.e.  $E = -126$ ,  $b_0 = 0$  ??

YES

## Relative Rounding Error, $x \neq 0$

The **relative rounding error** is defined by  $|\delta|$ ,

$$\delta \equiv \frac{\text{round}(x) - x}{x}.$$

Assuming  $x$  is in the normal range,

$$x = \pm m \times 2^E, \quad \text{where } m \geq 1,$$

so  $|x| \geq 2^E$ .

Since  $|\text{round}(x) - x| < \epsilon \times 2^E$ , for **all** rounding modes,

$$|\delta| < \frac{\epsilon \times 2^E}{2^E} = \epsilon.$$

## Relative Rounding Error, $x \neq 0$

The relative rounding error is bounded:

$$|\delta| < \epsilon$$

**Q:** Does the bound necessarily hold if  $0 < |x| < N_{\min}$ ,  
i.e.  $E = -126$  and  $b_0 = 0$  ?? Why ??

**NO.**

Since  $\delta = \frac{\text{round}(x)}{x} - 1$ , for any real  $x$  in the **normal range**,

$$\boxed{\text{round}(x) = x(1 + \delta), \quad |\delta| < \epsilon}$$

## An Important Idea

$$\text{round}(x) = x(1 + \delta),$$

so the **rounded value** of an arbitrary number  $x$  in the **normal range** is **equal to**  $x(1 + \delta)$ , where, regardless of the rounding mode,

$$|\delta| < \epsilon.$$

This is very important, because you can think of the stored value of  $x$  as **not exact**, but as **exact within a factor** of  $1 + \epsilon$ .

IEEE **single format numbers** are good to a factor of about  $1 + 10^{-7}$ , i.e., they **have about 7 accurate decimal digits**.

## Special Case of Round to Nearest

For **round to nearest**, the **absolute** rounding error can be **no more than half the gap between  $x_-$  and  $x_+$** .

In IEEE single, for all  $|x| = (b_0.b_1b_2\dots)_2 \times 2^E \leq N_{\max}$ ,

$$|\text{round}(x) - x| \leq 2^{-24} \times 2^E,$$

and in general

$$|\text{round}(x) - x| \leq \frac{1}{2}\epsilon \times 2^E.$$

For  $x$  in the **normal range** (so  $b_0 = 1$ )

$$\text{round}(x) = x(1 + \delta), \quad |\delta| \leq \frac{\frac{1}{2}\epsilon \times 2^E}{2^E} = \frac{1}{2}\epsilon.$$

## Recap

- IEEE floating point representation (Overton's Chap 4)

Format	Sign	Exponent	Fraction
Single (binary32)	1	8	23
double (binary64)	1	11	52
quadruple (binary128)	1	15	112

**Concepts:** hidden bit, exponent bias, machine precision, machine epsilon

- Rounding (Overton's Chap 5)

Four rounding modes: down, up, towards zero, to nearest

$$\frac{|\text{round}(x) - x|}{|x|} \begin{cases} < \epsilon, & \text{all rounding modes} \\ \leq \frac{1}{2}\epsilon, & \text{round to nearest} \end{cases}$$

where  $N_{\min} \leq |x| \leq N_{\max}$ .



## Today's topics

- Floating point operations (Overton's Chap 6)
- Exceptional situations (Overton's Chap 7)
- Floating point in C (Overton's Chap 10)

# Operations on Floating Point Numbers

The IEEE standard requires correctly rounded operations:

- correctly rounded basic arithmetic operations (+, −, \*, /);
- correctly rounded remainder and square root operations;
- correctly rounded format conversions.

**Correctly rounded** means rounded to fit the destination of the result, using rounding mode in effect.

## IEEE Rule for Rounding

The exact result of an operation may **not** be a floating point number, e.g. the **multiplication** of two **24-bit** significands generally gives a **48-bit** significand.

The IEEE standard requires that the computed result be the **correctly** rounded value of the **exact** result

## IEEE Rule for a Floating Point Operation

Let  $x$  and  $y$  be floating point numbers, and let  $\oplus, \ominus, \otimes, \oslash$  denote the **implementations** of  $+, -, *, /$  on the computer.

The **IEEE rule** for the basic arithmetic operations is then precisely:

$$x \oplus y = \text{round}(x + y),$$

$$x \ominus y = \text{round}(x - y),$$

$$x \otimes y = \text{round}(x \times y),$$

$$x \oslash y = \text{round}(x/y).$$

Therefore when  $x + y$  is in the **normal range**,

$$x \oplus y = (x + y)(1 + \delta), \quad |\delta| < \epsilon$$

for **all** rounding modes. Similarly for  $\ominus, \otimes$  and  $\oslash$ .

(Note that  $|\delta| \leq \epsilon/2$  for **round to nearest**).

## Format Conversion

The IEEE standard requires support for correctly rounded format conversions:

- Conversion between floating point numbers.
- Conversion between floating point and integer formats.
- Rounding a floating point number to an integral value (not an integer format).
- Binary to decimal and decimal to binary conversion.


## Exceptional Situations

When a reasonable response to exceptional data is possible, it should be used.

**Division by zero**. Two **earlier** standard responses :

- Generate the **largest FPN** as the result.  
Rationale: user would notice the large number in the output and conclude something had gone wrong.  
**Disaster**: e.g.  $2/0 - 1/0$  would then have a result of 0, which is **completely meaningless**. In general the user might **not even notice** that any error had taken place.
- Generate a **program interrupt**, e.g.  
**“fatal error — division by zero”**.  
The burden was on the programmer to make sure that division by zero would **never** occur.

## Example: computing the **total resistance**



The **total resistance**  $T = \frac{1}{\frac{1}{R_1} + \frac{1}{R_2}}$ .

What if  $R_1 = 0$ ? If one resistor offers no resistance, **all** the current will flow through that and avoid the other; therefore, the total resistance in the circuit is **zero**.

The formula for  $T$  also makes perfect sense **mathematically**:

$$T = \frac{1}{\frac{1}{0} + \frac{1}{R_2}} = \frac{1}{\infty + \frac{1}{R_2}} = \frac{1}{\infty} = 0.$$

## The IEEE FPS Solution

Why should a **programmer** have to worry about treating division by zero as an exceptional situation here?

In **IEEE floating point arithmetic**, **division by zero** does not generate an interrupt, but **gives an infinite result**, program execution continuing normally.

In the case of **the parallel resistance formula** this leads to a final **correct result** of  $1/\infty = 0$ , following the mathematical concepts **exactly**:

$$T = \frac{1}{\frac{1}{0} + \frac{1}{R_2}} = \frac{1}{\infty + \frac{1}{R_2}} = \frac{1}{\infty} = 0.$$

## Other uses of $\infty$

We used some of the following:

$$\begin{array}{ll} a > 0 & : \quad a/0 \rightarrow \infty, \quad a * \infty \rightarrow \infty, \\ a \text{ finite} & : \quad a + \infty \rightarrow \infty, \quad a - \infty \rightarrow -\infty, \\ & \quad a/\infty \rightarrow 0, \quad \infty + \infty \rightarrow \infty. \end{array}$$



## Uses of NaN

- The operations  $\infty * 0$ ,  $0/0$ ,  $\infty/\infty$ ,  $\infty - \infty$  are indefinite. Computing any of these is called an **invalid operation**, and the IEEE standard sets the result to NaN.
- A real operation with a complex result, e.g., the square root of a negative number, produces NaN.
- **Almost all** arithmetic operations with at least one NaN operand **also** produce NaN.
- Whenever a NaN is discovered in the output, the programmer **knows something has gone wrong**. An  $\infty$  in the output may or may not indicate an error, depending on the context.
- sNaN generates interruption while qNaN does not. The application decides if it generates qNaN or sNaN. For instance, GCC C compiler always generates qNaN unless explicitly specified to behave the other way around.

## Overflow and Underflow

**Overflow** is said to occur when

$$N_{\max} < | \text{true result} | < \infty,$$

where  $N_{\max}$  is the **largest** normal FPN.

Two **pre-IEEE** standard treatments:

- (i) Set the result to  $(\pm) N_{\max}$ , or
- (ii) Interrupt with an **error message**.

In IEEE arithmetic, the standard response depends on the **rounding mode**.

## Overflow and Underflow, ctd

Suppose that the overflowed value is **positive**. Then

rounding mode	result
<b>round up</b>	$\infty$
<b>round down</b>	$N_{\max}$
<b>round towards zero</b>	$N_{\max}$
<b>round to nearest</b>	$\infty$

**Round to nearest** is the **default** rounding mode and any other choice may lead to very misleading final computational results.

## Overflow and Underflow, ctd

- **Underflow** is said to occur when

$$0 < | \text{true result} | < N_{\min},$$

where  $N_{\min}$  is the **smallest** normal floating point number.

- Historically the response was usually:  
**replace the result by zero.**
- In **IEEE standard**, the result may be a **subnormal** number instead of zero – allowing results **much smaller** than  $N_{\min}$ .
- But there may still be a significant loss of accuracy, since subnormal numbers have fewer bits of format.

## Floating Point in C

In C, the type **float** refers to a **single format** FP variable.

*Example.* **read** in a FPN, using the standard input routine `scanf`, and **print** it out again, using `printf`:

```
main ()      /* echo.c:  echo the input */
{
    float x;
    scanf("%f", &x);
    printf("x = %f", x);
}
```

The 2nd argument `&x` to `scanf` is the **address** of `x`.

The routine `scanf` needs to know **where** to store the value read.

The 2nd argument `x` to `printf` is the **value** of `x`.

The 1st argument `"%f"` to both routines is a **control string**.

## Using Different Output Formats in “printf”

The following result was for a Sun 4.

With input 0.66666666666666666666 :

```
#cc -o echoinput echo.c
```

```
#echoinput
```

```
0.66666666666666666666          (typed in)
```

```
x= 0.666667                     (printed out)
```

## Format Codes

The two standard **format codes** used for specifying floating point numbers in these control strings are:

- `%f`, for **fixed decimal format**
- `%e`, for **exponential decimal format**

The two formats have **identical** effects in `scanf`, which can process input in a **fixed** decimal format (e.g. 0.666) or an **exponential** decimal format (e.g. 6.66e-1), but have different effects in `printf`.

## Using Different Output Formats in “printf”

Output format	Output
%f	0.666667
%e	6.666667e-01
%8.3f	0.667
%8.3e	6.667e-01
%20.15f	0.666666686534882
%20.15e	6.666666865348816e-01

- The input is rounded to about 6 or 7 digits of precision, so %f and %e print, **by default**, 6 digits after the decimal point.
- The next two lines print to **less** precision.
- In the last two lines about half the digits have **no significance**. Regardless of the **output format**, the floating point variables are **always** stored in the **IEEE formats**.



## Format Conversions: scanf and printf

- The scanf routine calls a decimal to binary conversion routine to convert the input decimal format to internal binary floating point representation;
- The printf routine calls a binary to decimal conversion routine to convert the binary floating point representation to the output decimal format.
- Both conversion routines use the rounding mode that is in effect to correctly round the results.

## A note on the %f format code

Using the %f format code is **NOT** a good idea, unless it is known the numbers are neither too small nor too large.  
e.g., if the input is  $1.0e-10$ , the output using %f is 0.000000.

In numerical computing, usually the %e format code is used.

## Double or Long Float

Double precision variables are declared in C using **double** or **long float**. But changing **float** to **double** in the previous program:

```
main ()      /* echo.c:  echo the input */
{ double x;
  scanf("%f",&x);
  printf("%e",x);
}
```

gives  $-6.392091e-236$ . **Q:** Why ??

`scanf` reads the input and stores it in **single format** in the first half of the **double word** allocated to `x`, but when `x` is **printed**, its value is read assuming it is **stored in double format**.

## Double or Long Float, ctd

- When `scanf` reads a **double precision** variable **we must use the format** `%lf` (for long float), so that it stores the result in **double** format.
- `printf` **expects** double precision, and single format variables are **automatically** converted to double before being passed to it.

Since it always receives **long float** arguments, it treats `%e` and `%le` identically;  
likewise `%f` and `%lf`, `%g` and `%lg`.

## Story: Effect of output format on a parliament election

Parliamentary election in Schleswig-Holstein, Germany, April 5, 1992.

- In the elections, a party with less than 5.0% of the vote cannot be seated.
- It was announced the Greens had a cliff-hanging 5.0% the evening of the election.
- It was discovered after midnight that they really had only 4.97%. The printout had one digit after the decimal point, and the actual percentage was rounded to 5.0%.

## A program to “test” if $x$ is “zero”

```
main() /* loop1.c: generate small numbers*/
{ float x; int n;
  n = 0; x = 1;    /* x = 2^0 */
  while (x != 0){
    n++;
    x = x/2;      /* x = 2^(-n) */
    printf("\n n= %d  x=%e", n,x); }
}
```

Initializes  $x$  to 1 and repeatedly divides by 2 until it rounds to 0.

## A program to “test” if $x$ is “zero”, ctd

Thus  $x$  becomes  $1/2, 1/4, 1/8, \dots$ , through **subnormal**  $2^{-127}$ ,  $2^{-128}$ , ... to **the smallest subnormal**  $2^{-149}$ .

The last value is 0, since  $2^{-150}$  is **not** representable.

**Output (various machines with various compliers):**

```
n= 1  x=5.000000e-01
n= 2  x=2.500000e-01
. . .
n= 149  x=1.401298e-45
n= 150  x=0.000000e+00
```

## Another “test” if $x$ is “zero”

```
main() /* loop2.c: this loop stops sooner*/
{ float x,y; int n;
  n = 0; y = 2; x = 1;      /* x = 2^0 */
  while (y != 1) {
    n++;
    x = x/2;      /* x = 2^(-n) */
    y = 1 + x;    /* y = 1 + 2^(-n) */
    printf("\n n= %d x= %e y= %e",n,x,y);
  }
}
```

**Initializes**  $x$  to 1 and repeatedly divides by 2, **terminating** when  $y = 1 + x$  is 1.



## Another “test” if $x$ is “zero”

This occurs **much sooner**, since  $1 + 2^{-24}$  is **not** a floating point number, rounds to 1. Note  $1 + 2^{-23}$  does **not** round to 1.

**Output (various machines with various compliers):**

```
n= 1 x= 5.000000e-01 y= 1.500000e+00
```

```
. . .
```

```
n= 23 x= 1.192093e-07 y= 1.000000e+00
```

```
n= 24 x= 5.960464e-08 y= 1.000000e+00
```

## Yet another “test” if $x$ is “zero”

Now instead of using the variable  $y$ , change the **test**

`while (y != 1) to while (1 + x != 1) :`

```
n=0;  x=1;      /* x = 2^0 */
```

```
while(1 + x != 1){... /*same as before*/}
```

It stops at

- $n = 24$  on a PC using Visual Studio or online GDB, and Mac.  
It uses registers with **the single precision**  $p = 24$
- $n = 53$  on a PC using Visual C++  
It uses registers with **the double precision**  $p = 24$
- $n = 64$  on a PC using gcc  
It uses registers with **the extended precision**  $p = 64$ .

**Conclusion:** Running the same program on different machines with different compilers may still give different results